

Task 1 — Convolutional Neural Networks (CNN) in Cybersecurity

Course: AI and ML for Cybersecurity

Student: Sofio Katamadze

1) What is a Convolutional Neural Network?

A Convolutional Neural Network (CNN) is a neural architecture specialized for learning local patterns in grid-like data. CNNs use **convolutional layers**—small learnable filters (kernels) that slide across the input—to detect features such as edges, textures, and shapes. Unlike fully connected layers, convolutions **reuse weights spatially**, drastically reducing parameters and improving sample efficiency. After convolutions, **nonlinear activations** (e.g., ReLU) introduce expressiveness, while **pooling** layers (e.g., max-pooling) downsample feature maps to gain translation invariance and reduce computation. Stacking multiple convolution–activation–pooling blocks yields hierarchical representations: early layers capture simple patterns; deeper layers compose them into higher-level concepts.

In cybersecurity, many signals can be reshaped as grids so CNNs can exploit locality: network-flow statistics arranged into small “images,” byte/entropy maps of files, or spectrograms of traffic timing. CNNs can learn subtle spatial correlations—e.g., co-occurring spikes in packet counts and flag combinations—that are hard to engineer by hand. Because convolutions are efficient and parallelizable on GPUs, CNNs scale to large datasets and near-real-time scenarios (e.g., intrusion detection on aggregated flows).

This task demonstrates a compact CNN that classifies network flows as **benign vs malicious** after mapping each flow’s 64 numeric features to an **8×8 grayscale image**. The pipeline: generate a synthetic flow dataset with a reproducible rule for attacks, normalize/reshape features into images, train a small CNN, and evaluate on held-out data. This image view lets the CNN discover localized feature groups that correlate with attacks, illustrating how spatial inductive bias improves learning even when the original signal is tabular.

2) Data Used

Preview table (first 8 rows):

==== DATA PREVIEW (first 8 rows) ===

```
f1  f2  f3  f4  f5  f6  f7  f8  f9  f10 f11 f12 f13 f14 f15 f16 f17 f18 f19 f20  
f21 f22 f23 f24 f25 f26 f27 f28 f29 f30 f31 f32 f33 f34 f35 f36 f37 f38 f39  
f40 f41 f42 f43 f44 f45 f46 f47 f48 f49 f50 f51 f52 f53 f54 f55 f56 f57 f58  
f59 f60 f61 f62 f63 f64 label
```

```
2.092 2.835 1.837 1.645 3.079 1.753 2.346 2.189 3.077 1.439 3.796 1.173 2.453 2.258 1.089  
2.595 1.524 0.813 2.822 0.566 0.947 1.822 3.064 2.706 2.068 0.404 1.130 1.336 0.780 0.301  
1.886 2.764 1.256 3.038 0.503 0.739 1.857 1.174 2.612 2.329 1.241 0.550 1.131 2.368 2.682  
1.178 0.300 0.476 2.238 1.225 1.248 0.729 0.716 3.004 2.293 1.011 1.762 1.985 4.708 0.664  
1.967 2.994 4.359 0.970 0
```

```
1.212 2.628 0.398 2.105 5.802 2.261 0.151 0.822 0.995 3.456 1.470 0.305 1.598 1.841 1.102  
0.705 6.201 3.000 3.182 1.473 0.956 1.146 0.501 4.649 1.518 3.985 1.190 2.284 1.915 0.451  
5.029 1.219 0.607 2.643 1.660 2.142 1.786 0.174 0.798 0.771 3.151 1.707 2.380 1.278 2.462  
1.475 1.394 1.905 1.786 7.487 0.790 0.401 2.108 0.891 0.147 0.637 6.461 0.772 0.272 6.204  
0.614 8.934 1.235 5.043 0
```

... (6 more rows omitted for brevity)

CSV snippet (first 12 rows):

==== CSV SNIPPET (first 12 rows) ===

```
f1,f2,f3,...,f64,label
```

```
2.0918174,2.8353455,1.8372155,...,0.9696665,0
```

```
1.2123756,2.6284661,0.39767373,...,5.0427446,0
```

... (9 more rows)

The dataset is generated by the script in Section 3 and then re-read from an in-memory CSV to satisfy the “data included” requirement while keeping the submission self-contained.

3) Python Code (data generation + reading + CNN)

```
task_1/task1_cnn_data_and_cnn.py
```

```
#!/usr/bin/env python3
```

```
"""
```

```
task1_cnn_data_and_cnn.py — Self-contained CNN demo with DATA GENERATED + READ  
in-code.
```

```
"""
```

```
import io  
  
import numpy as np  
  
import pandas as pd  
  
import tensorflow as tf  
  
from tensorflow.keras import layers, models
```

```
def generate_flow_dataset(n_samples: int = 200, n_features: int = 64, seed: int = 42):  
  
    rng = np.random.default_rng(seed)  
  
    X = rng.gamma(shape=2.0, scale=1.0, size=(n_samples, n_features)).astype("float32")  
  
    hot_idx = np.array([1, 3, 5, 7, 11, 13, 17, 19, 23, 29])  
  
    score = X[:, hot_idx].sum(axis=1)  
  
    y = (score > np.percentile(score, 65)).astype("int32")  
  
    return X, y
```

```
def dataframe_with_labels(X: np.ndarray, y: np.ndarray) -> pd.DataFrame:  
  
    cols = [f"f{i+1}" for i in range(X.shape[1])]  
  
    df = pd.DataFrame(X, columns=cols)  
  
    df["label"] = y  
  
    return df
```

```
def to_csv_text(df: pd.DataFrame, max_rows: int = None) -> str:  
  
    if max_rows is not None:
```

```
df = df.head(max_rows)

buf = io.StringIO()

df.to_csv(buf, index=False)

return buf.getvalue()

def read_from_csv_text(csv_text: str) -> pd.DataFrame:

    return pd.read_csv(io.StringIO(csv_text))

def build_cnn(input_shape=(8, 8, 1)):

    model = models.Sequential([
        layers.Input(shape=input_shape),
        layers.Conv2D(16, (3, 3), activation="relu"),
        layers.MaxPool2D(2, 2),
        layers.Conv2D(32, (3, 3), activation="relu"),
        layers.Flatten(),
        layers.Dense(32, activation="relu"),
        layers.Dense(1, activation="sigmoid"),
    ])

    model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy"])

    return model

def main():

    X, y = generate_flow_dataset(n_samples=200, n_features=64, seed=42)

    df = dataframe_with_labels(X, y)
```

```
print("\n==== DATA PREVIEW (first 8 rows) ====")

print(df.head(8).round(3).to_string(index=False))

csv_preview = to_csv_text(df, max_rows=12)

print("\n==== CSV SNIPPET (first 12 rows) ====")

print(csv_preview)

_ = read_from_csv_text(csv_preview) # demonstration of "reading the data"

X_full = X.reshape(-1, 8, 8, 1)

y_full = y

idx = np.arange(len(y_full))

rng = np.random.default_rng(123)

rng.shuffle(idx)

n_train = int(0.8 * len(idx))

train_idx, test_idx = idx[:n_train], idx[n_train:]

Xtr, Xte = X_full[train_idx], X_full[test_idx]

ytr, yte = y_full[train_idx], y_full[test_idx]

model = build_cnn(input_shape=(8, 8, 1))

_ = model.fit(Xtr, ytr, epochs=10, batch_size=16, validation_split=0.2, verbose=0)
```

```
loss, acc = model.evaluate(Xte, yte, verbose=0)

print(f"\nTest Accuracy = {acc:.3f} (loss={loss:.3f})")

if __name__ == "__main__":
    import os

    os.environ["TF_CPP_MIN_LOG_LEVEL"] = "2"

    main()
```

4) Results

- **Test Accuracy:** 0.625 (loss 0.672) from the run
- Notes:
 - 64 per-flow numeric features are mapped to **8×8** grayscale images.
 - Labels mark flows with high mass on selected “hot” features as **malicious** (toy rule).
 - Data are generated and **re-read** within the same script—no external files required.

5) Reproduce

```
cd task_1

python3 task1_cnn_data_and_cnn.py
```